



# African Masters Of Machine Intelligence AMMI

## Explaining and harnessing adversarial examples

authors: Ian J. Goodfellow, Jonathon Shlens, Christian Szegedy

Presented by

ENGELBERT TCHINDE  
ewamba@aimsammi.org

Supervised by  
Tutors

August 27, 2021

1. Motivation
2. Adversarial example
3. Adversarial example for linear model
4. Linear perturbation for non-linear model

# Motivation

- Models are not learning the true underlying properties of the data.


# Motivation

- Models are not learning the true underlying properties of the data.
- The cause of adversarial examples is a mystery.

# Motivation

- Models are not learning the true underlying properties of the data.
- The cause of adversarial examples is a mystery.
  - Extreme nonlinearity of deep neural networks
  - Insufficient model averaging
  - Insufficient regularization

# Adversarial example



# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.

# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.
- Consider a linear model that take an input image  $\mathbf{x}$ .



# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.
- Consider a linear model that take an input image  $\mathbf{x}$ .
- Image: 8 bits for pixel value. 256 pixel values (0-255).

# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.
- Consider a linear model that take an input image  $\mathbf{x}$ .
- Image: 8 bits for pixel value. 256 pixel values (0-255).
- Perturbation  $\eta$  of each pixel of the image  $\mathbf{x}$ .
  - $\tilde{\mathbf{x}} = \mathbf{x} + \eta$ ,
  - $\|\eta\|_{\infty} < \epsilon$ .

# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.
- Consider a linear model that take an input image  $\mathbf{x}$ .
- Image: 8 bits for pixel value. 256 pixel values (0-255).
- Perturbation  $\eta$  of each pixel of the image  $\mathbf{x}$ .
  - $\tilde{\mathbf{x}} = \mathbf{x} + \eta$ ,
  - $\|\eta\|_{\infty} < \epsilon$ .
- Why minimization?
  - If we did not minimize  $\eta$ , we could just get a feature vector  $\mathbf{x}$ .

# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.
- Consider a linear model that take an input image  $\mathbf{x}$ .
- Image: 8 bits for pixel value. 256 pixel values (0-255).
- Perturbation  $\eta$  of each pixel of the image  $\mathbf{x}$ .
  - $\tilde{\mathbf{x}} = \mathbf{x} + \eta$ ,
  - $\|\eta\|_{\infty} < \epsilon$ .
- Why minimization?
  - If we did not minimize  $\eta$ , we could just get a feature vector  $\mathbf{x}$ .
  - We want to keep the semantic (meaning) of the initial picture by applying so many noise.

# Adversarial example for linear model

- In many problems, the precision of an individual input feature is limited.
- Consider a linear model that take an input image  $\mathbf{x}$ .
- Image: 8 bits for pixel value. 256 pixel values (0-255).
- Perturbation  $\eta$  of each pixel of the image  $\mathbf{x}$ .
  - $\tilde{\mathbf{x}} = \mathbf{x} + \eta$ ,
  - $\|\eta\|_{\infty} < \epsilon$ .
- Why minimization?
  - If we did not minimize  $\eta$ , we could just get a feature vector  $\mathbf{x}$ .
  - We want to keep the semantic (meaning) of the initial picture by applying so many noise.
- By considering the dot product between a weight vector  $w$  and an adversarial example  $\tilde{\mathbf{x}}$ ,

$$\mathbf{w}^T \tilde{\mathbf{x}} = \mathbf{w}^T \mathbf{x} + \mathbf{w}^T \eta.$$

# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \boldsymbol{\eta}$  ?

# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \boldsymbol{\eta}$  ?
  - The adversarial perturbation causes the activation to grow by  $\mathbf{w}^T \boldsymbol{\eta}$ .

# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \eta$  ?
  - The adversarial perturbation causes the activation to grow by  $\mathbf{w}^T \eta$ .
- We define the perturbation to be the sign of the weight vector  $\mathbf{w}$ . Why?



# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \eta$  ?
  - The adversarial perturbation causes the activation to grow by  $\mathbf{w}^T \eta$ .
- We define the perturbation to be the sign of the weight vector  $\mathbf{w}$ . Why?
  - We are trying to maximize this dot product  $\mathbf{w}^T \eta$  s.t.  $\|\eta\|_\infty < \epsilon$

# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \eta$  ?
  - The adversarial perturbation causes the activation to grow by  $\mathbf{w}^T \eta$ .
- We define the perturbation to be the sign of the weight vector  $\mathbf{w}$ . Why?
  - We are trying to maximize this dot product  $\mathbf{w}^T \eta$  s.t.  $\|\eta\|_\infty < \epsilon$
- The activation will grow by  $\epsilon m n$ ,

# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \boldsymbol{\eta}$  ?
  - The adversarial perturbation causes the activation to grow by  $\mathbf{w}^T \boldsymbol{\eta}$ .
- We define the perturbation to be the sign of the weight vector  $\mathbf{w}$ . Why?
  - We are trying to maximize this dot product  $\mathbf{w}^T \boldsymbol{\eta}$  s.t.  $\|\boldsymbol{\eta}\|_{\infty} < \epsilon$
- The activation will grow by  $\epsilon mn$ ,
  - $n = \dim \mathbf{w}$
  - $m =$  average magnitude of an element of  $\mathbf{w}$ .

# Adversarial example for linear model

- Why we control  $\mathbf{w}^T \eta$  ?
  - The adversarial perturbation causes the activation to grow by  $\mathbf{w}^T \eta$ .
- We define the perturbation to be the sign of the weight vector  $\mathbf{w}$ . Why?
  - We are trying to maximize this dot product  $\mathbf{w}^T \eta$  s.t.  $\|\eta\|_\infty < \epsilon$
- The activation will grow by  $\epsilon mn$ ,
  - $n = \dim \mathbf{w}$
  - $m =$  average magnitude of an element of  $\mathbf{w}$ .
- *A simple linear model can have adversarial examples if its input has sufficient dimensionality.*

# Linear perturbation for non-linear model



$x$


# Linear perturbation for non-linear model



$x$

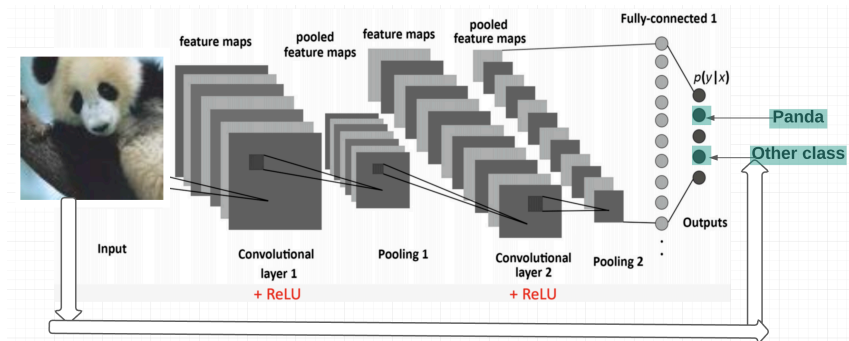
- Select a random real world image  $x$ : **panda**.

# Linear perturbation for non-linear model



- Run the input image  $\mathbf{x}$  into a ConvNet and get a correct classifier as a **panda**.

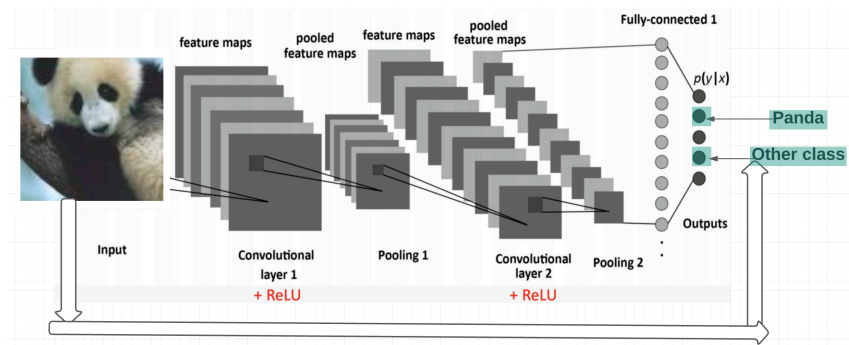
# Linear perturbation for non-linear model



- Run the input image  $\mathbf{x}$  into a ConvNet and get a correct classifier as a **panda**.
- Select a random output neuron in the output layer that is different from the true neuron that classifies **panda**.




# Linear perturbation for non-linear model



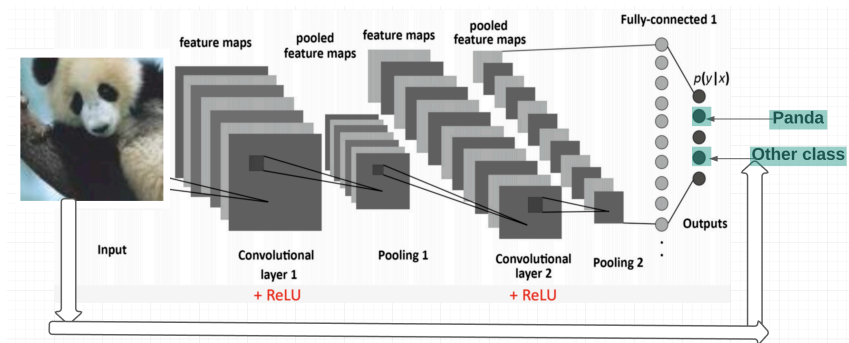
- We apply **GD** to the input pixels of our **panda** in order to minimize the classification loss with respect to the newly neuron chosen class neuron.

# Linear perturbation for non-linear model



- We apply **GD** to the input pixels of our **panda** in order to minimize the classification loss with respect to the newly neuron chosen class neuron.
- Instead of adjusting the network weights in order to optimize our classifier, we adjust the input pixels until they fool the network to make a wrong prediction.

# Linear perturbation for non-linear model



- The final trick is to make sure that our generated image looks as close as possible to the original one such that we can't see the difference.


# Linear perturbation for non-linear model

- We can generate the perturbation  $\eta$  using a Fast Gradient Sign Methode (FGSM),
  - $\eta = \epsilon * \text{sign}(\nabla_x J(\theta, \mathbf{x}, y))$ ,
  - $\theta$  the parameter,  $\mathbf{x}$  input vector,  $y$  target associated to  $\mathbf{x}$ ,  $J$  cost function.


 $\mathbf{x}$ 

“panda”

57.7% confidence

 $+ .007 \times$ 

 $\text{sign}(\nabla_x J(\theta, \mathbf{x}, y))$ 

“nematode”

8.2% confidence

 $=$ 

 $\mathbf{x} +$ 
 $\epsilon \text{sign}(\nabla_x J(\theta, \mathbf{x}, y))$ 

“gibbon”

99.3 % confidence

# Experiments

From paper:

<b>epsilon</b>	<b>Error rate</b>	<b>Confidence</b>	<b>Activation</b>	<b>Dataset</b>
0.25	99.9%	79.3%	Softmax	MNIST
0.25	89.4%	97.6%	Maxout	MNIST
0.10	87.15%	96.6%	Maxout	CIFAR-10

# Experiments

From paper:

<b>epsilon</b>	<b>Error rate</b>	<b>Confidence</b>	<b>Activation</b>	<b>Dataset</b>
0.25	99.9%	79.3%	Softmax	MNIST
0.25	89.4%	97.6%	Maxout	MNIST
0.10	87.15%	96.6%	Maxout	CIFAR-10

Our experiments:

<b>epsilon</b>	<b>Error rate</b>	<b>Confidence</b>	<b>Activation</b>	<b>Dataset</b>
0.25	67.33%	0.29%	LogSoftmax	MNIST
0.1 0	15.36%	0.39%	LogSoftmax	MNIST

<b>epsilon</b>	<b>Error rate</b>	<b>Confidence</b>	<b>Activation</b>	<b>Dataset</b>
0.25	63.79%	99.78%	Softmax	MNIST
0.1 0	14.07%	99.64%	Softmax	MNIST

# End

## THANK YOU !